

# ECN 142 Final Project

Jayden Kwong

2025-03-06

## Import Libraries

```
library(leaps)
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-8

library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(corrplot)




## corrplot 0.92 loaded

data <- read.csv("telecom_churn.csv")
data <- na.omit(data)
set.seed(123)


# Set graphical layout to display multiple plots
par(mfrow = c(2, 2)) # Adjust rows and columns based on the number of predictors

# List of predictors to plot
predictors <- c("AccountWeeks", "ContractRenewal", "DataPlan", "DataUsage",
                 "CustServCalls", "DayMins", "DayCalls", "MonthlyCharge",
                 "OverageFee", "RoamMins")


# Loop through each predictor and create a boxplot
for (var in predictors) {
  boxplot(data[[var]] ~ data$Churn,
          main = paste("Boxplot of", var, "by Churn"),
          xlab = "Churn (0 = No, 1 = Yes)",
          ylab = var,
          col = c("blue", "red"))
}
```

**Boxplot of AccountWeeks by Churn****Boxplot of ContractRenewal by Churn****Boxplot of DataPlan by Churn**


Churn (0 = No, 1 = Yes)

**Boxplot of DataUsage by Churn**


Churn (0 = No, 1 = Yes)

**Boxplot of CustServCalls by Churn**


Churn (0 = No, 1 = Yes)

**Boxplot of DayMins by Churn**

Churn (0 = No, 1 = Yes)


**Boxplot of DayCalls by Churn**

Churn (0 = No, 1 = Yes)

**Boxplot of MonthlyCharge by Churn**

Churn (0 = No, 1 = Yes)


**Boxplot of OverageFee by Churn**



Churn (0 = No, 1 = Yes)

Logistic Regression with all Model

**Boxplot of RoamMins by Churn**



Churn (0 = No, 1 = Yes)

#

```
set.seed(123)

# 80% Train
trainIndex <- createDataPartition(data$Churn, p = 0.8, list = FALSE)
trainData <- data[trainIndex, ]
testData <- data[-trainIndex, ]

log_model <- glm(Churn ~ ., data = trainData, family = "binomial")
summary(log_model)
```

```
##
## Call:
## glm(formula = Churn ~ ., family = "binomial", data = trainData)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -5.6349292  0.6080743 -9.267 < 2e-16 ***
## AccountWeeks   0.0006289  0.0015459   0.407  0.68415
## ContractRenewal -1.9652691  0.1578396 -12.451 < 2e-16 ***
## DataPlan      -1.4420265  0.5977300 -2.413  0.01584 *
## DataUsage       0.9285898  2.1373236   0.434  0.66395
## CustServCalls   0.5118968  0.0423880  12.076 < 2e-16 ***
## DayMins        0.0247914  0.0360776   0.687  0.49198
## DayCalls        0.0026905  0.0030612   0.879  0.37946
## MonthlyCharge   -0.0755197  0.2119522  -0.356  0.72161
## OverageFee       0.2616866  0.3617678   0.723  0.46946
## RoamMins        0.0772774  0.0244615   3.159  0.00158 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2201.8 on 2666 degrees of freedom
## Residual deviance: 1748.0 on 2656 degrees of freedom
## AIC: 1770
##
## Number of Fisher Scoring iterations: 6
```

```
test_pred <- predict(log_model, newdata = testData, type = "response")
test_class <- ifelse(test_pred > 0.5, 1, 0)
```


```
accuracy <- mean(test_class == testData$Churn)
cat("Logistic Regression Accuracy:", accuracy, "\n")
```

## Logistic Regression Accuracy: 0.8618619

## Look at Correlation Matrix

### Polynomial Logistic Regression

```
#cor(data)
corrplot(cor(data))
```



```
# Load necessary libraries
library(splines)

# Fit a logistic regression with polynomial terms for selected variables
log_poly <- glm(Churn ~ poly(CustServCalls, 2) + poly(OverageFee, 2) + poly(MonthlyCharge, 2) +
                  ContractRenewal + DataUsage,
                  family = "binomial", data = trainData)

summary(log_poly)
```

```

## 
## Call:
## glm(formula = Churn ~ poly(CustServCalls, 2) + poly(OverageFee,
##      2) + poly(MonthlyCharge, 2) + ContractRenewal + DataUsage,
##      family = "binomial", data = trainData)
##
## Coefficients:
##                               Estimate Std. Error z value Pr(>|z|)
## (Intercept)                 0.4849    0.1547   3.134  0.00172 ***
## poly(CustServCalls, 2)1 30.9098    2.9200  10.586 < 2e-16 ***
## poly(CustServCalls, 2)2 20.9084    3.4608   6.042 1.53e-09 ***
## poly(OverageFee, 2)1     0.1209    3.6381   0.033  0.97349
## poly(OverageFee, 2)2     0.5283    3.2640   0.162  0.87142
## poly(MonthlyCharge, 2)1 66.6549    6.1890  10.770 < 2e-16 ***
## poly(MonthlyCharge, 2)2 12.6253    3.9133   3.226  0.00125 **
## ContractRenewal          -1.9564    0.1560 -12.540 < 2e-16 ***
## DataUsage                  -1.1941    0.1140 -10.478 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2201.8  on 2666  degrees of freedom
## Residual deviance: 1724.0  on 2658  degrees of freedom
## AIC: 1742
##
## Number of Fisher Scoring iterations: 6

# Predict probabilities on the test set
pred_probs <- predict(log_poly, newdata = testData, type = "response")

# Convert probabilities into binary predictions (Threshold = 0.5)
pred_labels <- ifelse(pred_probs > 0.5, 1, 0)

# Compute Accuracy
accuracy <- mean(pred_labels == testData$Churn)
print(paste("Model Accuracy:", round(accuracy, 4)))

## [1] "Model Accuracy: 0.8649"

```

## Variable Selection (Best Subset Selection)

```

best_subset <- regsubsets(Churn ~ ., data= data, nvmax = 10)
subset_summary <- summary(best_subset)

best_adj <- which.max(subset_summary$adjr2)
best_cp <- which.min(subset_summary$cp)
best_bic <- which.min(subset_summary$bic)

best_model_adj <- coef(best_subset, best_adj)
best_model_cp <- coef(best_subset, best_cp)

```

```

best_model_bic <- coef(best_subset, best_bic)

print(best_model_adj)

##      (Intercept) ContractRenewal      DataUsage CustServCalls      DayCalls
## -0.1470461398   -0.3001137523   -0.1029598144    0.0582735286  0.0003467457
## MonthlyCharge           RoamMins
## 0.0074675218    0.0098361679

print(best_model_cp)

##      (Intercept) ContractRenewal      DataUsage CustServCalls MonthlyCharge
## -0.112451288   -0.300175125   -0.103006680    0.058173176  0.007465555
##          RoamMins
## 0.009894002

print(best_model_bic) # Use BIC

##      (Intercept) ContractRenewal      DataUsage CustServCalls MonthlyCharge
## -0.112451288   -0.300175125   -0.103006680    0.058173176  0.007465555
##          RoamMins
## 0.009894002

```

## Variable Selection (Forward Step Selection)

```

regfit.fwd <- regsubsets(Churn ~ ., data = data, nvmax = 10, method = "forward")
fwd_summary <- summary(regfit.fwd)

fwd_adj <- which.max(fwd_summary$adjr2)
fwd_cp <- which.min(fwd_summary$cp)
fwd_bic <- which.min(fwd_summary$bic)

fwd_model_adj <- coef(regfit.fwd, fwd_adj)
fwd_model_cp <- coef(regfit.fwd, fwd_cp)
fwd_model_bic <- coef(regfit.fwd, fwd_bic)

print(fwd_model_adj)

##      (Intercept) ContractRenewal      DataPlan CustServCalls      DayMins
## -0.1244205134   -0.2992595976   -0.0795901811    0.0583184330  0.0012624977
##          DayCalls      OverageFee           RoamMins
## 0.0003485952    0.0128758547    0.0077235678

print(fwd_model_cp)

##      (Intercept) ContractRenewal      DataPlan CustServCalls      DayMins
## -0.089255626   -0.299317791   -0.079761287    0.058217071  0.001263361
##      OverageFee           RoamMins
## 0.012817028   0.007776409


```

```
print(fwd_model_bic) # Use BIC
```

```
##      (Intercept) ContractRenewal      DataPlan   CustServCalls      DayMins
## -0.089255626    -0.299317791    -0.079761287     0.058217071     0.001263361
##      OverageFee       RoamMins
##     0.012817028     0.007776409
```

## Variable Selection (LASSO)

```
library(glmnet)
x <- model.matrix(Churn ~ ., data = data)[,-1]
y <- data$Churn
lasso_cv <- cv.glmnet(x, y, alpha = 1, family = "binomial")
plot(lasso_cv)
```



```
best_lambda <- lasso_cv$lambda.min
print(best_lambda)
```

```
## [1] 0.001525839
```

```
lasso_coef <- coef(lasso_cv, s = "lambda.min")
print(lasso_coef)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##           s1
```

```

## (Intercept) -5.6357616768
## AccountWeeks 0.0002844391
## ContractRenewal -1.9394322642
## DataPlan -0.8764699438
## DataUsage .
## CustServCalls 0.4907927477
## DayMins 0.0123204950
## DayCalls 0.0028228982
## MonthlyCharge .
## OverageFee 0.1306506269
## RoamMins 0.0761396580

lasso_predictions <- predict(lasso_cv, newx = x, s = "lambda.min", type = "response")
pred_class <- ifelse(lasso_predictions > 0.5, 1, 0)
lasso_accuracy <- mean(pred_class == y)
lasso_accuracy

## [1] 0.860486

```

## Comparing each Models

```

set.seed(123)

# Split into training (80%) and test (20%)
trainIndex <- createDataPartition(data$Churn, p = 0.8, list = FALSE)
trainData <- data[trainIndex, ]
testData <- data[-trainIndex, ]

# Train Logistic Regression using Best Subset Features
selected_features_subset <- names(coef(best_subset, best_bic))[-1]
log_model_subset <- glm(Churn ~ ., data = trainData[, c("Churn", selected_features_subset)], family = "binomial")

# Train Logistic Regression using Forward Selection Features
selected_features_fwd <- names(coef(best_subset, fwd_bic))[-1]
log_model_fwd <- glm(Churn ~ ., data = trainData[, c("Churn", selected_features_fwd)], family = "binomial")

# Train Logistic Regression using LASSO Features
selected_features_lasso <- rownames(lasso_coef)[lasso_coef[,1] != 0][-1] # Remove Intercept
log_model_lasso <- glm(Churn ~ ., data = trainData[, c("Churn", selected_features_lasso)], family = "binomial")

# Make Predictions on Test Data
pred_subset <- predict(log_model_subset, newdata = testData, type = "response")
pred_fwd <- predict(log_model_fwd, newdata = testData, type = "response")
pred_lasso <- predict(log_model_lasso, newdata = testData, type = "response")

# Convert Probabilities to Binary (0 or 1)
pred_class_subset <- ifelse(pred_subset > 0.5, 1, 0)
pred_class_fwd <- ifelse(pred_fwd > 0.5, 1, 0)
pred_class_lasso <- ifelse(pred_lasso > 0.5, 1, 0)

```

```

# Compute Accuracy
acc_subset <- mean(pred_class_subset == testData$Churn)
acc_fwd <- mean(pred_class_fwd == testData$Churn)
acc_lasso <- mean(pred_class_lasso == testData$Churn)

# Accuracy Results
cat("Best Subset Selection Accuracy:", acc_subset, "\n")

## Best Subset Selection Accuracy: 0.8618619

cat("Forward Selection Accuracy:", acc_fwd, "\n")

## Forward Selection Accuracy: 0.8648649

cat("LASSO Accuracy:", acc_lasso, "\n")

## LASSO Accuracy: 0.8618619

regfit.fwd <- regsubsets(Churn ~ ., data = data, nvmax = 10, method = "forward")
fwd_summary <- summary(regfit.fwd)

best_model_size_fwd <- which.min(fwd_summary$bic)
selected_features_fwd <- names(coef(regfit.fwd, best_model_size_fwd))[-1]

# Selected Features
cat("Stepwise Selected Features:\n")

## Stepwise Selected Features:

print(selected_features_fwd)

## [1] "ContractRenewal" "DataPlan"           "CustServCalls"    "DayMins"
## [5] "OverageFee"      "RoamMins"

# Prepare predictor and response
x <- model.matrix(Churn ~ ., data = data[, c("Churn", selected_features_fwd)])[, -1]
y <- as.numeric(as.character(data$Churn))

# Perform LASSO with Cross-Validation
set.seed(123)
lasso_cv <- cv.glmnet(x, y, alpha = 1, family = "binomial")

# Best lambda
best_lambda <- lasso_cv$lambda.min
cat("Best Lambda for LASSO:", best_lambda, "\n")

## Best Lambda for LASSO: 0.0007248971

```

```

# Coefficients after choosing lambda
lasso_coef <- coef(lasso_cv, s = "lambda.min")
print(lasso_coef)

## 7 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -5.43959997
## ContractRenewal -1.96541674
## DataPlan      -0.90674248
## CustServCalls 0.49790446
## DayMins       0.01256298
## OverageFee    0.13438030
## RoamMins     0.08010583

selected_features_lasso <- rownames(lasso_coef)[lasso_coef[,1] != 0][-1]

# Final Model
final_model_lasso <- glm(Churn ~ ., data = data[, c("Churn", selected_features_lasso)], family = "binomial")

summary(final_model_lasso)

##
## Call:
## glm(formula = Churn ~ ., family = "binomial", data = data[, c("Churn",
##   selected_features_lasso)])
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -5.552897  0.432757 -12.831 < 2e-16 ***
## ContractRenewal -1.989219  0.143452 -13.867 < 2e-16 ***
## DataPlan      -0.934814  0.144015 -6.491 8.52e-11 ***
## CustServCalls  0.505651  0.038834  13.021 < 2e-16 ***
## DayMins       0.012774  0.001073  11.907 < 2e-16 ***
## OverageFee    0.138612  0.022648   6.120 9.34e-10 ***
## RoamMins      0.083476  0.020304   4.111 3.93e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2758.3 on 3332 degrees of freedom
## Residual deviance: 2190.6 on 3326 degrees of freedom
## AIC: 2204.6
##
## Number of Fisher Scoring iterations: 5

# Evaluate Accuracy
# Split into training & test sets
set.seed(123)
trainIndex <- createDataPartition(data$Churn, p = 0.8, list = FALSE)
trainData <- data[trainIndex, ]
testData <- data[-trainIndex, ]

```

```

# Train model on training data
final_model_lasso <- glm(Churn ~ ., data = trainData[, c("Churn", selected_features_lasso)], family = "binomial")

# Predict on test data
test_pred <- predict(final_model_lasso, newdata = testData, type = "response")

# Convert probabilities to binary (0 or 1)
test_class <- ifelse(test_pred > 0.5, 1, 0)

# Compute Accuracy
final_accuracy <- mean(test_class == testData$Churn)
cat("Final Model Accuracy After Stepwise + LASSO:", final_accuracy, "\n")

```


## Final Model Accuracy After Stepwise + LASSO: 0.8633634

## Check for Multicollinearity

```

library(corrplot)
numeric_features <- trainData[, selected_features_lasso]
numeric_features <- numeric_features[, sapply(numeric_features, is.numeric)]
corr_matrix <- cor(numeric_features, use = "complete.obs")
corrplot(corr_matrix, method = "color", type = "upper",
          tl.cex = 0.8, number.cex = 0.7)

```



## PCA

```
# Standardize features before PCA (important for variance scaling)
preProc <- preProcess(trainData[, selected_features_lasso], method = "pca", pcaComp = 5)

# Transform data using PCA
train_pca <- predict(preProc, trainData[, selected_features_lasso])
test_pca <- predict(preProc, testData[, selected_features_lasso])

# Add the target variable back
train_pca$Churn <- trainData$Churn
test_pca$Churn <- testData$Churn

# Fit logistic regression on PCA-transformed data
pca_log_model <- glm(Churn ~ ., data = train_pca, family = "binomial")

# Summary of PCA-based logistic model
summary(pca_log_model)

##
## Call:
## glm(formula = Churn ~ ., family = "binomial", data = train_pca)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.25277   0.07629 -29.529 < 2e-16 ***
## PC1         -0.57872   0.05423 -10.671 < 2e-16 ***
## PC2          0.54435   0.06445  8.446 < 2e-16 ***
## PC3          0.06816   0.06127  1.113   0.266
## PC4          0.86837   0.06090  14.258 < 2e-16 ***
## PC5          0.34574   0.06773  5.105  3.31e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2201.8 on 2666 degrees of freedom
## Residual deviance: 1753.2 on 2661 degrees of freedom
## AIC: 1765.2
##
## Number of Fisher Scoring iterations: 5

# Predict probabilities
pca_pred_probs <- predict(pca_log_model, newdata = test_pca, type = "response")

# Convert to class predictions (Threshold = 0.5)
pca_pred_labels <- ifelse(pca_pred_probs > 0.5, 1, 0)

# Compute accuracy
pca_accuracy <- mean(pca_pred_labels == test_pca$Churn)
print(paste("PCA Logistic Regression Accuracy:", round(pca_accuracy, 4)))

## [1] "PCA Logistic Regression Accuracy: 0.8634"
```

## Classification Tree

```
library(tree)
library(rpart)
library(rpart.plot)
library(caret) # For model evaluation

set.seed(42)
train_index <- sample(1:nrow(trainData), size = 0.7 * nrow(trainData))
valid_index <- setdiff(1:nrow(trainData), train_index)

# Define training and validation datasets
train <- trainData[train_index, c("Churn", selected_features_lasso)]
valid <- trainData[valid_index, c("Churn", selected_features_lasso)]

# Train the Classification Tree
tree_model <- rpart(Churn ~ ., data = train, method = "class")


printcp(tree_model)

##
## Classification tree:
## rpart(formula = Churn ~ ., data = train, method = "class")
##
## Variables actually used in tree construction:
## [1] ContractRenewal CustServCalls DataPlan      DayMins
## [5] OverageFee     RoamMins
##
## Root node error: 271/1866 = 0.14523
##
## n= 1866
##
##          CP nsplit rel error  xerror      xstd
## 1 0.097786      0    1.00000 1.00000 0.056162
## 2 0.051661      2    0.80443 0.80443 0.051201
## 3 0.040590      5    0.64945 0.75646 0.049847
## 4 0.036900      7    0.56827 0.65683 0.046824
## 5 0.033210      8    0.53137 0.65683 0.046824
## 6 0.014760      9    0.49815 0.57196 0.043991
## 7 0.011070     11    0.46863 0.56458 0.043732
## 8 0.010000     13    0.44649 0.54244 0.042941

optimal_cp <- tree_model$cptable[which.min(tree_model$cptable[, "xerror"]), "CP"]
print(paste("Optimal CP:", optimal_cp))

## [1] "Optimal CP: 0.01"

pruned_tree <- prune(tree_model, cp = optimal_cp)
rpart.plot(pruned_tree, type = 3, fallen.leaves = TRUE, cex = 0.8)
```



```

conf_matrix <- confusionMatrix(predict(pruned_tree, newdata = testData, type = "class"), as.factor(test$Class))
print(conf_matrix)

```

```

## Confusion Matrix and Statistics
##
##          Reference
## Prediction 0 1
##          0 555 37
##          1 13 61
##
##          Accuracy : 0.9249
##          95% CI : (0.9022, 0.9438)
##          No Information Rate : 0.8529
##          P-Value [Acc > NIR] : 8.83e-09
##
##          Kappa : 0.6672
##
##  Mcnemar's Test P-Value : 0.001143
##
##          Sensitivity : 0.9771
##          Specificity : 0.6224
##          Pos Pred Value : 0.9375
##          Neg Pred Value : 0.8243
##          Prevalence : 0.8529
##          Detection Rate : 0.8333
##          Detection Prevalence : 0.8889
##          Balanced Accuracy : 0.7998
##
```

```
##      'Positive' Class : 0  
##
```